

# 优化方法一

12700 8核16线程, 80gb内存 文件放在share memory上 写入到nvme固态

## 优化一：多进程并行计算（MPI）

相关库：mpi4py

优化说明：

初始版本完全串行处理所有图像帧且不区分任务分工；最终版本引入 MPI，将图像帧的处理任务在多个进程中并行执行，并通过 rank 和 size 明确分配数据处理范围，最后聚合处理结果。

初始版本代码片段1：

```
def reslicing(self,name):
    """
    用刚体变换参数将图片进行重采样对齐\n
    Resampling and aligning the image with the rigid-body transform
    parameter\n
    output: 对齐后的图片(Aligned image)
    """
    self.interpolator =
itk.BSplineInterpolateImageFunction.New(self.img_gauss) # B样条插值器(B-
spline interpolation)

    print("开始重采样对齐(注意, 需要很长时间, 建议先去干点别的事)\nreslicing")
    new = np.ndarray(self.shape)
    for picture in range(self.shape[0]):
        new[picture,:, :, :] =
self.get_new_img(self.img_data_gauss[picture,:, :, :],
self.parameter[picture],picture)
        print(f"进度{picture*100/self.shape[0]}%", end="\r")
    self.save_img(new, name)
    return new
```

最终代码版本1

```
unit = self.shape[0] // size
remain = self.shape[0] % size
start = rank * unit + remain if rank != 0 else 0
end = (rank + 1) * unit + remain

seperate_img = np.empty((end - start, ) + tuple(self.shape[1:]))
for picture in range(start, end):
    seperate_img[picture - start, :, :, :] = self.get_new_img(
        self.img_data_gauss[picture,:, :, :],
```

```

        self.parameter[picture],
        picture
    )

if rank == 0:
    gathered_img = np.empty((unit * size, ) + tuple(self.shape[1:]))
    comm.Gather(seperate_img[remain:], gathered_img, root=0)
    self.save_img(gathered_img, name)
else:
    comm.Gather(seperate_img, None, root=0)
python
unit = self.shape[0] // size
remain = self.shape[0] % size
start = rank * unit + remain if rank != 0 else 0
end = (rank + 1) * unit + remain

seperate_img = np.empty((end - start, ...))
for picture in range(start, end):
    seperate_img[picture-start, :, :, :] = self.get_new_img(...)

comm.Gather(seperate_img, gathered_img, root=0)

```

初始版本代码片段2(iterate 函数中的三维点分布):

```

def iterate(self, reference, resource, q, pic_num):
    """
    高斯牛顿迭代(gauss-newton iterate)\n
    输入参数(parameter):\n
    resource:原始图像(raw image)\n
    reference:参考图像(reference image)\n
    q:旋转平移参数(rotation and translation parameters)\n
    输出(output):\n
    q:更新后的旋转平移参数(new rotation and translation parameters)\n
    bi:残差(residual)\n
    b_diff:偏导矩阵(derivative matrix)
    """

    interpolator = self.interpolator # B样条插值(B-spline
interpolation)
    step = self.step # 选点间隔(point interval)
    bi = np.zeros(self.x*self.y*self.z) # 残差 (residual)
    # 偏导矩阵 (derivative matrix)
    b_diff = np.zeros(((self.x*self.y*self.z), 7))
    index = 0
    for i in range(self.x):
        for j in range(self.y):
            for k in range(self.z):
                n_i = i*step # 对应位置的转换,获得点在原图的位置(get the
position of the point in the original image)
                n_j = j*step
                n_k = k*step

```

```

M = self.rigid(q)
M[3,3]=0.0
if np.linalg.det(M)==0:
    interpo_pos = np.linalg.pinv(M)[n_i, n_j, n_k, 1]
else:
    interpo_pos =np.linalg.inv(M)[n_i, n_j, n_k, 1]
point = itk.Point[itk.D, 4]()
if 0 <= interpo_pos[0] < self.shape[1] and 0 <=
interpo_pos[1] < self.shape[2] and 0 <= interpo_pos[2] < self.shape[3]: #
判断是否在范围内

    point[0] = pic_num
    point[1] = interpo_pos[0]
    point[2] = interpo_pos[1]
    point[3] = interpo_pos[2]

else:
    point[1] = n_i
    point[2] = n_j
    point[3] = n_k
    point[0] = pic_num

bi[index] = (interpolator.Evaluate(point)-
reference[n_i][n_j][n_k])**2
tem=interpolator.Evaluate(point)-reference[n_i][n_j]
[n_k]

derivative = interpolator.EvaluateDerivative(point)
diff_x = derivative[1]
diff_y = derivative[2]
diff_z = derivative[3]
(a,b,c)=self.shape[1:]
b_diff[index][1:4] = diff_x*tem, diff_y*tem, diff_z*tem
b_diff[index][4] = (diff_y*(-0.5*self.affine[0][0]*a*
(sin(q[4])*sin(q[6])) - sin(q[5])*cos(q[4])*cos(q[6]))/self.affine[1][1] +
self.affine[0][0]*(sin(q[4])*sin(q[6])) -
sin(q[5])*cos(q[4])*cos(q[6]))/self.affine[1][1] - 0.5*b*(-
sin(q[4])*cos(q[6]) - sin(q[5])*sin(q[6])*cos(q[4])) - sin(q[4])*cos(q[6])
- sin(q[5])*sin(q[6])*cos(q[4]) - 0.5*self.affine[2]
[2]*c*cos(q[4])*cos(q[5])/self.affine[1][1] + self.affine[2]
[2]*cos(q[4])*cos(q[5])/self.affine[1][1]) + \
diff_z*(-0.5*self.affine[0][0]*a*
(sin(q[4])*sin(q[5])*cos(q[6]) + sin(q[6])*cos(q[4]))/self.affine[2][2] +
self.affine[0][0]*(sin(q[4])*sin(q[5])*cos(q[6]) +
sin(q[6])*cos(q[4]))/self.affine[2][2] - 0.5*self.affine[1][1]*b*(
sin(q[4])*sin(q[5])*sin(q[6]) -
cos(q[4])*cos(q[6]))/self.affine[2][2] + self.affine[1][1]*
(sin(q[4])*sin(q[5])*sin(q[6]) - cos(q[4])*cos(q[6]))/self.affine[2][2] +
0.5*c*sin(q[4])*cos(q[5]) - sin(q[4])*cos(q[5])))*tem

b_diff[index][5] = (diff_x*(0.5*a*sin(q[5])*cos(q[6])) -
sin(q[5])*cos(q[6]) + 0.5*self.affine[1]
[1]*b*sin(q[5])*sin(q[6])/self.affine[0][0] - self.affine[1]
[1]*sin(q[5])*sin(q[6])/self.affine[0][0] - 0.5*self.affine[2]
[2]*c*cos(q[5])/self.affine[0][0] + self.affine[2]
[2]*cos(q[5])/self.affine[0][0]) +\

```

```

        diff_y*(0.5*self.affine[0]
[0]*a*sin(q[4])*cos(q[5])*cos(q[6])/self.affine[1][1] - self.affine[0]
[0]*sin(q[4])*cos(q[5])*cos(q[6])/self.affine[1][1] +
0.5*b*sin(q[4])*sin(q[6])*cos(q[5]) - sin(q[4])*sin(q[6])*cos(q[5]) +
0.5*self.affine[2][2]*c*sin(q[4])*sin(q[5])/self.affine[1][1] -
self.affine[2][2]*sin(q[4])*sin(q[5])/self.affine[1][1]) +\
        diff_z*(0.5*self.affine[0]
[0]*a*cos(q[4])*cos(q[5])*cos(q[6])/self.affine[2][2] - self.affine[0]
[0]*cos(q[4])*cos(q[5])*cos(q[6])/self.affine[2][2] + 0.5*self.affine[1]
[1]*b*sin(
        q[6])*cos(q[4])*cos(q[5])/self.affine[2][2] -
self.affine[1][1]*sin(q[6])*cos(q[4])*cos(q[5])/self.affine[2][2] +
0.5*c*sin(q[5])*cos(q[4]) - sin(q[5])*cos(q[4]))*tem

        b_diff[index][6] = (diff_x*(0.5*a*sin(q[6])*cos(q[5]) -
sin(q[6])*cos(q[5]) - 0.5*self.affine[1]
[1]*b*cos(q[5])*cos(q[6])/self.affine[0][0] + self.affine[1]
[1]*cos(q[5])*cos(q[6])/self.affine[0][0]) +\
        diff_y*(-0.5*self.affine[0][0]*a*
(sin(q[4])*sin(q[5])*sin(q[6]) - cos(q[4])*cos(q[6]))/self.affine[1][1] +
self.affine[0][0]*(sin(q[4])*sin(q[5])*sin(q[6]) -
cos(q[4])*cos(q[6]))/self.affine[1][1] - 0.5*b*(-
sin(q[4])*sin(q[5])*cos(q[6]) - sin(q[6])*cos(q[4])) -
sin(q[4])*sin(q[5])*cos(q[6]) - sin(q[6])*cos(q[4])) +\
        diff_z*(-0.5*self.affine[0][0]*a*
(sin(q[4])*cos(q[6]) + sin(q[5])*sin(q[6])*cos(q[4]))/self.affine[2][2] +
self.affine[0][0]*(sin(q[4])*cos(q[6]) +
sin(q[5])*sin(q[6])*cos(q[4]))/self.affine[2]
[2] - 0.5*self.affine[1][1]*b*
(sin(q[4])*sin(q[6]) - sin(q[5])*cos(q[4])*cos(q[6]))/self.affine[2][2] +
self.affine[1][1]*(sin(q[4])*sin(q[6]) -
sin(q[5])*cos(q[4])*cos(q[6]))/self.affine[2][2]))*tem

        b_diff[index][0] = 1
        index += 1
    return bi, b_diff

```

## 最终代码版本2

```

def iterate(self, reference, resource, q, pic_num):
    """
    高斯牛顿迭代(gauss-newton iterate)\n
    输入参数(parameter):\n
    resource:原始图像(raw image)\n
    reference:参考图像(reference image)\n
    q:旋转平移参数(rotation and translation parameters)\n
    输出(output):\n
    q:更新后的旋转平移参数(new rotation and translation parameters)\n
    bi:残差(residual)\n
    b_diff:偏导矩阵(derivative matrix)
    """

```

```

        interpolator = self.interpolator # B样条插值(B-spline
interpolation)
        step = self.step # 选点间隔(point interval)

        total = self.x * self.y * self.z
        unit = total // size # 每个进程处理的图像数量
        remain = total % size # 剩余的图像数量

        if rank == 0:
            start = 0
        else:
            start = rank * unit + remain

        end = (rank + 1) * unit + remain

        bi = np.zeros(end - start) # 残差 (residual)
        b_diff = np.zeros((end - start, 7)) # 偏导矩阵 (derivative matrix)

        M = self.rigid(q)
        M[3,3]=0.0
        if np.linalg.det(M) <= 1e-10:
            M = np.linalg.pinv(M)
        else:
            M = np.linalg.inv(M)

        # 处理每个进程分配的图像
        for index in range(start, end):
            i = index // (self.y * self.z)
            j = (index // self.z) % self.y
            k = index % self.z

            n_i = int(i*step * self.affine[1, 1]) if int(i*step *
self.affine[1, 1]) < self.shape[1] else self.shape[1] - 1
            n_j = int(j*step * self.affine[2, 2]) if int(j*step *
self.affine[2, 2]) < self.shape[2] else self.shape[2] - 1
            n_k = int(k*step * self.affine[3, 3]) if int(k*step *
self.affine[3, 3]) < self.shape[3] else self.shape[3] - 1

            interpo_pos = M @ [n_i, n_j, n_k, 1] # 变换后的坐标(transformed
coordinate)

            point = itk.Point[itk.D, 4]()
            if 0 <= interpo_pos[0] < self.shape[1] and 0 <= interpo_pos[1]
< self.shape[2] and 0 <= interpo_pos[2] < self.shape[3]: # 判断是否在范围内
                point[0] = pic_num
                point[1] = interpo_pos[0]
                point[2] = interpo_pos[1]
                point[3] = interpo_pos[2]
            else:
                point[1] = n_i
                point[2] = n_j
                point[3] = n_k

```

```

        point[0] = pic_num
        # bi[index - start] = (interpolator.Evaluate(point)-
reference[n_i][n_j][n_k])**2
        bi[index - start] = self.use_inter_evaluate(point,
interpolator, reference, n_i, n_j, n_k)
        # tem=interpolator.Evaluate(point)-reference[n_i][n_j][n_k]
        # derivative = interpolator.EvaluateDerivative(point)
        derivative = self.use_inter_evaluate_derivative(point,
interpolator)
        diff_x = derivative[1]
        diff_y = derivative[2]
        diff_z = derivative[3]
        (a,b,c)=self.shape[1:]
        b_diff[index - start][1:4] = diff_x, diff_y, diff_z
        # b_diff[index - start][1:4] = derivative[1:4]
        b_diff[index - start][4] = (diff_y*(-0.5*self.affine[0][0]*a*
(sin(q[4])*sin(q[6])) - sin(q[5])*cos(q[4])*cos(q[6]))/self.affine[1][1] +
self.affine[0][0]*(sin(q[4])*sin(q[6])) -
sin(q[5])*cos(q[4])*cos(q[6]))/self.affine[1][1] - 0.5*b*(-
sin(q[4])*cos(q[6]) - sin(q[5])*sin(q[6])*cos(q[4])) - sin(q[4])*cos(q[6])
- sin(q[5])*sin(q[6])*cos(q[4]) - 0.5*self.affine[2]
[2]*c*cos(q[4])*cos(q[5])/self.affine[1][1] + self.affine[2]
[2]*cos(q[4])*cos(q[5])/self.affine[1][1]) + \
                diff_z*(-0.5*self.affine[0][0]*a*
(sin(q[4])*sin(q[5])*cos(q[6]) + sin(q[6])*cos(q[4]))/self.affine[2][2] +
self.affine[0][0]*(sin(q[4])*sin(q[5])*cos(q[6]) +
sin(q[6])*cos(q[4]))/self.affine[2][2] - 0.5*self.affine[1][1]*b*(
                sin(q[4])*sin(q[5])*sin(q[6]) -
cos(q[4])*cos(q[6]))/self.affine[2][2] + self.affine[1][1]*
(sin(q[4])*sin(q[5])*sin(q[6]) - cos(q[4])*cos(q[6]))/self.affine[2][2] +
0.5*c*sin(q[4])*cos(q[5]) - sin(q[4])*cos(q[5])) #*tem
        b_diff[index - start][5] = (diff_x*(0.5*a*sin(q[5])*cos(q[6])) -
sin(q[5])*cos(q[6]) + 0.5*self.affine[1]
[1]*b*sin(q[5])*sin(q[6])/self.affine[0][0] - self.affine[1]
[1]*sin(q[5])*sin(q[6])/self.affine[0][0] - 0.5*self.affine[2]
[2]*c*cos(q[5])/self.affine[0][0] + self.affine[2]
[2]*cos(q[5])/self.affine[0][0]) +\
                diff_y*(0.5*self.affine[0]
[0]*a*sin(q[4])*cos(q[5])*cos(q[6])/self.affine[1][1] - self.affine[0]
[0]*sin(q[4])*cos(q[5])*cos(q[6])/self.affine[1][1] +
0.5*b*sin(q[4])*sin(q[6])*cos(q[5]) - sin(q[4])*sin(q[6])*cos(q[5]) +
0.5*self.affine[2][2]*c*sin(q[4])*sin(q[5])/self.affine[1][1] -
self.affine[2][2]*sin(q[4])*sin(q[5])/self.affine[1][1]) +\
                diff_z*(0.5*self.affine[0]
[0]*a*cos(q[4])*cos(q[5])*cos(q[6])/self.affine[2][2] - self.affine[0]
[0]*cos(q[4])*cos(q[5])*cos(q[6])/self.affine[2][2] + 0.5*self.affine[1]
[1]*b*sin(
                q[6])*cos(q[4])*cos(q[5])/self.affine[2][2] -
self.affine[1][1]*sin(q[6])*cos(q[4])*cos(q[5])/self.affine[2][2] +
0.5*c*sin(q[5])*cos(q[4]) - sin(q[5])*cos(q[4]))#*tem
        b_diff[index - start][6] = (diff_x*(0.5*a*sin(q[6])*cos(q[5]) -
sin(q[6])*cos(q[5]) - 0.5*self.affine[1]
[1]*b*cos(q[5])*cos(q[6])/self.affine[0][0] + self.affine[1]
[1]*cos(q[5])*cos(q[6])/self.affine[0][0]) +\

```

```

        diff_y*(-0.5*self.affine[0][0]*a*
(sin(q[4])*sin(q[5])*sin(q[6]) - cos(q[4])*cos(q[6]))/self.affine[1][1] +
self.affine[0][0]*(sin(q[4])*sin(q[5])*sin(q[6]) -
cos(q[4])*cos(q[6]))/self.affine[1][1] - 0.5*b*(-
sin(q[4])*sin(q[5])*cos(q[6]) - sin(q[6])*cos(q[4])) -
sin(q[4])*sin(q[5])*cos(q[6]) - sin(q[6])*cos(q[4])) +\
        diff_z*(-0.5*self.affine[0][0]*a*(sin(q[4])*cos(q[6]) +
sin(q[5])*sin(q[6])*cos(q[4]))/self.affine[2][2] + self.affine[0][0]*
(sin(q[4])*cos(q[6]) + sin(q[5])*sin(q[6])*cos(q[4]))/self.affine[2][2] -
0.5*self.affine[1][1]*b*(sin(q[4])*sin(q[6]) -
sin(q[5])*cos(q[4])*cos(q[6]))/self.affine[2][2] + self.affine[1][1]*
(sin(q[4])*sin(q[6]) - sin(q[5])*cos(q[4])*cos(q[6]))/self.affine[2]
[2]))#*tem

        b_diff[index - start][0] = 1

        # print('tem:',tem)
        # break

# 汇总所有进程处理的结果
gathered_bi = np.zeros(unit * size) # 创建一个空数组用于存储汇总结果
gathered_b_diff = np.zeros((unit * size, 7)) # 创建一个空数组用于存储
汇总结果

        if rank == 0:
            comm.Gather(bi[remain:], gathered_bi, root=0) # 将每个进程处理的
            结果汇总到主进程
            comm.Gather(b_diff[remain:], gathered_b_diff, root=0) # 将每个
            进程处理的结果汇总到主进程
        else:
            comm.Gather(bi, gathered_bi, root=0) # 将每个进程处理的结果汇总到
            主进程
            comm.Gather(b_diff, gathered_b_diff, root=0) # 将每个进程处理的结
            果汇总到主进程
        return gathered_bi, gathered_b_diff

```

## 加速原理

使用 mpi4py 将计算分发到多个进程（可分布于多核心或多节点），有效利用 CPU 多核资源。

明确分工避免不同进程重复处理数据。

每个进程只创建和使用自身所需内存，降低整体内存占用。

主进程统一收集结果，便于后续统一处理和保存

## 加速结果

1200s -> 200s

## 优化二：矩阵逆计算复用

优化说明：

刚体变换矩阵  $M$  每次在像素级调用 `np.linalg.inv()` 会极大拖慢计算。最终版本在外部计算一次 `invM` 后传入内部使用。

### 初始版本代码片段

```
for i in range(...):
    M = self.rigid(q)
    interpo_pos = np.linalg.inv(M) @ [i, j, k, 1]
```

### 最终版本

```
M = self.rigid(q)
invM = np.linalg.inv(M) if np.linalg.det(M) > 1e-10 else np.linalg.pinv(M)
...
for i in prange(...):
    pos = invM @ np.array([i, j, k, 1])
```

### 加速原理

矩阵逆是一项昂贵操作，将其提至循环外只算一次，大幅减少冗余计算。

### 加速结果

200s -> 40s

## 优化三：JIT 加速与并行循环（Numba）

相关库： `numba`

### 优化说明：

三重 for-loop 是图像处理中的瓶颈，初始版本直接使用原生 Python 实现，执行效率极低。最终版本使用 `@jit(parallel=True)` 修饰函数，编译为高性能本地代码，并开启线程级并行。

### 初始版本代码片段：

```
def get_new_img(...):
    for i in range(self.shape[1]):
        for j in range(self.shape[2]):
            for k in range(self.shape[3]):
                interpo_pos = ...
                new_img[i, j, k] = ...
```

### 最终版本



```
@njit(parallel=True)
def fast_get_new_img(resource, invM, shape1, shape2, shape3):
    new_img = np.empty((shape1, shape2, shape3))
    for i in prange(shape1):
        for j in range(shape2):
            for k in range(shape3):
                pos = invM @ np.array([i, j, k, 1.0])
                ...
```

## 加速原理

@njit: 使用 Numba 的 JIT 编译器将 Python 函数编译为本地代码。prange: 开启多线程并行循环，加快大规模像素点处理速度。

## 加速结果

40s -> 30s

## 优化四 MPI 进程绑定策略（核心绑定）

涉及指令：mpirun --bind-to hwthread --map-by core

## 优化说明：

在多核 CPU 系统中，MPI 进程如果调度不当，可能在不同 CPU 核之间迁移，导致缓存频繁失效，影响性能。使用 --bind-to hwthread 或 --map-by core 等参数可以将每个 MPI 进程固定到指定的逻辑核心或物理核心上，提升数据局部性和执行效率。

## 加速结果

30s -> 20s